

```
In [1]: # K-Nearest Neighbors (K-NN)

# Importing the libraries
import os
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
In [2]: os.chdir("D:/My ML Simulations\\My_ML_Work/Part 3 - Classification/Section 15 - K-Nearest Neighbors (K-NN)")
```

```
In [3]: # Importing the dataset
dataset = pd.read_csv('Social_Network_Ads.csv')
X = dataset.iloc[:, [2, 3]].values
y = dataset.iloc[:, 4].values
```

```
In [14]: dataset.head(9)
```

Out[14]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
5	15728773	Male	27	58000	0
6	15598044	Female	27	84000	0
7	15694829	Female	32	150000	1
8	15600575	Male	25	33000	0

```
In [15]: # Splitting the dataset into the Training set and Test set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 0)
```

```
In [16]: # Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [35]: print(X_test, 'sc.fit_transform')
```

[[-0.80480212 0.50496393]
[-0.01254409 -0.5677824]
[-0.30964085 0.1570462]
[-0.80480212 0.27301877]
[-0.30964085 -0.5677824]
[-1.10189888 -1.43757673]
[-0.70576986 -1.58254245]
[-0.21060859 2.15757314]
[-1.99318916 -0.04590581]
[0.8787462 -0.77073441]
[-0.80480212 -0.59677555]
[-1.00286662 -0.42281668]
[-0.11157634 -0.42281668]
[0.08648817 0.21503249]
[-1.79512465 0.47597078]
[-0.60673761 1.37475825]
[-0.11157634 0.21503249]
[-1.89415691 0.44697764]
[1.67100423 1.75166912]
[-0.30964085 -1.37959044]
[-0.30964085 -0.65476184]
[0.8787462 2.15757314]
[0.28455268 -0.53878926]
[0.8787462 1.02684052]
[-1.49802789 -1.20563157]
[1.07681071 2.07059371]
[-1.00286662 0.50496393]
[-0.90383437 0.30201192]
[-0.11157634 -0.21986468]
[-0.60673761 0.47597078]
[-1.6960924 0.53395707]
[-0.11157634 0.27301877]
[1.86906873 -0.27785096]
[-0.11157634 -0.48080297]
[-1.39899564 -0.33583725]
[-1.99318916 -0.50979612]
[-1.59706014 0.33100506]
[-0.4086731 -0.77073441]
[-0.70576986 -1.03167271]
[1.07681071 -0.97368642]
[-1.10189888 0.53395707]
[0.28455268 -0.50979612]
[-1.10189888 0.41798449]
[-0.30964085 -1.43757673]
[0.48261718 1.22979253]
[-1.10189888 -0.33583725]
[-0.11157634 0.30201192]
[1.37390747 0.59194336]
[-1.20093113 -1.14764529]
[1.07681071 0.47597078]
[1.86906873 1.51972397]
[-0.4086731 -1.29261101]

```
[ -0.30964085 -0.3648304 ]
[ -0.4086731  1.31677196]
[  2.06713324  0.53395707]
[  0.68068169 -1.089659  ]
[ -0.90383437  0.38899135]
[ -1.20093113  0.30201192]
[  1.07681071 -1.20563157]
[ -1.49802789 -1.43757673]
[ -0.60673761 -1.49556302]
[  2.1661655  -0.79972756]
[ -1.89415691  0.18603934]
[ -0.21060859  0.85288166]
[ -1.89415691 -1.26361786]
[  2.1661655  0.38899135]
[ -1.39899564  0.56295021]
[ -1.10189888 -0.33583725]
[  0.18552042 -0.65476184]
[  0.38358493  0.01208048]
[ -0.60673761  2.331532  ]
[ -0.30964085  0.21503249]
[ -1.59706014 -0.19087153]
[  0.68068169 -1.37959044]
[ -1.10189888  0.56295021]
[ -1.99318916  0.35999821]
[  0.38358493  0.27301877]
[  0.18552042 -0.27785096]
[  1.47293972 -1.03167271]
[  0.8787462   1.08482681]
[  1.96810099  2.15757314]
[  2.06713324  0.38899135]
[ -1.39899564 -0.42281668]
[ -1.20093113 -1.00267957]
[  1.96810099 -0.91570013]
[  0.38358493  0.30201192]
[  0.18552042  0.1570462  ]
[  2.06713324  1.75166912]
[  0.77971394 -0.8287207  ]
[  0.28455268 -0.27785096]
[  0.38358493 -0.16187839]
[ -0.11157634  2.21555943]
[ -1.49802789 -0.62576869]
[ -1.29996338 -1.06066585]
[ -1.39899564  0.41798449]
[ -1.10189888  0.76590222]
[ -1.49802789 -0.19087153]
[  0.97777845 -1.06066585]
[  0.97777845  0.59194336]
[  0.38358493  0.99784738]] sc.fit_transform
```

```
In [17]: # Fitting K-NN to the Training set
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors = 5, metric = 'minkowski', p = 2)
classifier.fit(X_train, y_train)
```

```
Out[17]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                             weights='uniform')
```

```
In [18]: # Predicting the Test set results
y_pred = classifier.predict(X_test)
```

```
In [31]: print(y_pred, 'classifier.predict')
```

```
[0 0 0 0 0 0 1 0 1 0 0 0 0 0 1 0 0 1 0 1 0 1 0 0 0 0 0 1 0 0 0 0
 0 0 1 0 0 0 0 1 0 0 1 0 1 1 0 0 1 1 1 0 0 1 0 0 1 0 1 0 1 0 0 0 1 0 0 1
 0 0 0 0 1 1 1 1 0 0 1 0 0 1 1 0 0 1 0 0 0 0 0 1 1 1] classifier.predict
```

```
In [19]: # Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
```


```
In [42]: print(cm, 'confusion_matrix')
```

```
[[64  4]
 [ 3 29]] confusion_matrix
```

```
In [24]: # Visualising the Training set results
from matplotlib.colors import ListedColormap
X_set, y_set = X_train, y_train
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
                     np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()]).T).reshape(X1.shape),
             alpha = 0.75, cmap = ListedColormap(('blue', 'red')))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('blue', 'red'))(i), label = j)
plt.title('K-NN (Training set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.



```
In [21]: # Visualising the Test set results
from matplotlib.colors import ListedColormap
X_set, y_set = X_test, y_test
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
                      np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))
plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()]).T).reshape(X1.shape),
             alpha = 0.75, cmap = ListedColormap(('red', 'green')))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c = ListedColormap(('red', 'green'))(i), label = j)
plt.title('K-NN (Test set)')
plt.xlabel('Age')
plt.ylabel('Estimated Salary')
plt.legend()
plt.show()
```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.

